

Book Recommendation System

CP421 Data Mining Project

Arsalan Khan

Asil Khan

Waleed Ghufraan

Spencer Robson


Aliha Ali

Qusai Almasri

Table of Contents

Content:


1. Introduction
2. Dataset Description
3. Data Preprocessing
4. Methodologies
5. RMSE Comparison
6. Discussion
7. Conclusion
8. Questions & Answers



Data Mining

['dā-tə 'mī-niŋ]

A process used by companies to turn raw data into useful information.

 Investopedia



Book Recommendation

Introduction

- Importance of recommendation systems in the information age.
- Project objectives:
 - Implement various recommendation approaches.
 - Compare their performance using evaluation metrics.
 - Enhance system accuracy with advanced methods.

Dataset Description

□ Book-Crossing Dataset:

- Users.csv: User demographics.
- Books.csv: Book metadata.
- Ratings.csv: User ratings.

□ Dataset Statistics:

- Users: 278,858 entries.
- Books: 271,360 entries.
- Ratings: 1,149,780 entries.
- Ratings scale: 1 to 10.

Users Summary:

	User-ID	Location	Age
count	278858.00000	278858	278858.000000
unique	NaN	57339	NaN
top	NaN	london, england, united kingdom	NaN
freq	NaN	2506	NaN
mean	139429.50000	NaN	34.754753
std	80499.51502	NaN	14.040587
min	1.00000	NaN	10.000000
25%	69715.25000	NaN	24.000000
50%	139429.50000	NaN	32.000000
75%	209143.75000	NaN	44.000000
max	278858.00000	NaN	100.000000

Books Summary:

	ISBN	Book-Title	Book-Author	Year-Of-Publication
count	271360	271360	271360	271360.0
unique	271360	242135	102022	116.0
top	0195153448	Selected Poems	Agatha Christie	2002.0
freq	1	27	632	17627.0

	Publisher	Image-URL-S \
count	271360	271360
unique	16807	271044
top	Harlequin	http://images.amazon.com/images/P/185326119X.0...
freq	7535	2

	Image-URL-M	Image-URL-L
count	271360	271360
unique	271044	271042
top	http://images.amazon.com/images/P/185326119X.0...	No Image
freq	2	3

Ratings Summary:

	User-ID	ISBN	Book-Rating
count	433671.000000	433671	433671.000000
unique	NaN	185973	NaN
top	NaN	0316666343	NaN
freq	NaN	707	NaN
mean	135458.743451	NaN	7.601066
std	80678.385078	NaN	1.843798
min	8.000000	NaN	1.000000
25%	66619.000000	NaN	7.000000
50%	133184.000000	NaN	8.000000
75%	205735.000000	NaN	9.000000
max	278854.000000	NaN	10.000000

Data Preprocessing

- Importance of data preprocessing.
- Steps undertaken:
 - Handling missing values.
 - Data filtering.
 - Data splitting.

Filtered Ratings: 96234
Training Set: 76987, Test Set: 19247

```
# Data Cleaning and Preprocessing
def filter_and_split_ratings(ratings, min_user_ratings=5, min_item_ratings=10, test_size=0.2, random_state=42):
    # Filter users with sufficient interactions
    user_counts = ratings['User-ID'].value_counts()
    valid_users = user_counts[user_counts >= min_user_ratings].index
    filtered_ratings = ratings[ratings['User-ID'].isin(valid_users)]


    # Filter items with sufficient interactions
    item_counts = filtered_ratings['ISBN'].value_counts()
    valid_items = item_counts[item_counts >= min_item_ratings].index
    filtered_ratings = filtered_ratings[filtered_ratings['ISBN'].isin(valid_items)]

    # Ensure sufficient data remains after filtering
    if len(filtered_ratings) < 100:
        raise ValueError("Filtered dataset is too small. Adjust filtering thresholds.")

    # Split data into train and test sets
    train_data, test_data = sklearn_train_test_split(filtered_ratings, test_size=test_size, random_state=random_state)

    # Print summary
    print(f"Filtered Ratings: {len(filtered_ratings)}")
    print(f"Training Set: {len(train_data)}, Test Set: {len(test_data)}")

    return train_data, test_data, filtered_ratings
```

Handling Missing Values

- **Users Dataset:**
 - Imputed missing ages by sampling valid ages.
 - Clipped ages to 10-100 years.
- **Books Dataset:**
 - 'Year-Of-Publication': Filled invalid years with 'Unknown'.
 - 'Book-Author' and 'Publisher': Filled with 'Unknown'.
 - 'Image-URL-L': Filled with 'No Image'.
- **Ratings Dataset:**
 - Removed ratings with a value of 0.

Data Filtering and Splitting

□ Data Filtering:

- Retained users with at least 5 ratings.
- Retained books with at least 10 ratings.

□ Data Splitting:

- 80/20 split for training and testing.
- Ensures models are evaluated on unseen data.

□ Resulting Data Sizes:

- Filtered Ratings: 96,234 entries.
- Training Set: 76,987 entries.
- Testing Set: 19,247 entries.



We implement three baseline methods:

1. 🌐 **Global Mean:**
 - The average rating across all books and users.
2. 👤 **User Mean:**
 - The average rating given by each user.
3. 📖 **Item Mean:**
 - The average rating received by each book.

These baselines provide simple reference points to compare the performance of more advanced models.

Collaborative Filtering Prediction (SVD)

try:

```
svd_prediction = svd_model.predict(user_id, item_id).est
```

except:

```
svd_prediction = ratings['Book-Rating'].mean()
```

Content-Based Filtering Prediction

```
book_index = books[books['ISBN'] == item_id].index
```

if not book_index.empty:

```
content_scores = content_sim_matrix[book_index[0]]
```

```
top_indices = content_scores.argsort()[-6:-1] # Get top 5 similar books
```

```
content_score = np.mean(content_scores[top_indices])
```

else:

```
content_score = 0.0
```

Methodology Overview

Overview of recommendation techniques implemented:

- Baseline Methods
- Content-Based Filtering
- Collaborative Filtering (SVD)
- Neural Collaborative Filtering (NCF)
- Hybrid Recommendation System

```
Epoch 1/20
151/151 ██████████ 5s 14ms/step - loss: 40.3619 - mae: 5.6047 - val_loss: 2.8695 - val_mae: 1.3587
Epoch 2/20
151/151 ██████████ 2s 8ms/step - loss: 4.6134 - mae: 1.7183 - val_loss: 2.6767 - val_mae: 1.3056
Epoch 3/20
151/151 ██████████ 1s 8ms/step - loss: 3.8251 - mae: 1.5514 - val_loss: 2.5743 - val_mae: 1.2608
Epoch 4/20
151/151 ██████████ 2s 8ms/step - loss: 3.4515 - mae: 1.4691 - val_loss: 2.5816 - val_mae: 1.2670
Epoch 5/20
151/151 ██████████ 1s 8ms/step - loss: 3.2146 - mae: 1.4206 - val_loss: 2.6077 - val_mae: 1.2734
Epoch 6/20
151/151 ██████████ 1s 8ms/step - loss: 3.0457 - mae: 1.3766 - val_loss: 2.6021 - val_mae: 1.2710
Epoch 6: early stopping
Restoring model weights from the end of the best epoch: 3.
602/602 ██████████ 1s 1ms/step - loss: 2.5322 - mae: 1.2487
NCF Test Loss (MSE): 2.5743191242218018
NCF Test Mean Absolute Error (MAE): 1.2607836723327637
```


Baseline Methods

- Purpose: Provide reference points for model performance.
- Methods:
 - Global Mean
 - User Mean
 - Item Mean

```
# Baseline Methods

# Global Mean Baseline
def calculate_global_mean(ratings):
    return ratings['Book-Rating'].mean()

# User Mean Baseline
def calculate_user_means(ratings):
    return ratings.groupby('User-ID')['Book-Rating'].mean()

# Item Mean Baseline
def calculate_item_means(ratings):
    return ratings.groupby('ISBN')['Book-Rating'].mean()

# Compute global mean
global_mean = calculate_global_mean(filtered_ratings)
print("Global Mean Rating:", global_mean)
```

Baseline Methods Results

- **Global Mean Rating: 7.8236**
- **User Mean RMSE: 1.6774**
- **Item Mean RMSE: 1.7476**

We implement three baseline methods:

1. 🌐 **Global Mean:**
 - The average rating across all books and users.
2. 👤 **User Mean:**
 - The average rating given by each user.
3. 📖 **Item Mean:**
 - The average rating received by each book.

These baselines provide simple reference points to compare the performance of more advanced models.

Content-Based Filtering

- ❑ Recommends items similar to those liked by the user.
- ❑ Based on item features.

Approach

- ❑ **TF-IDF Vectorization:**
 - Converts text into numerical features.
 - Captures importance of words in titles.
- ❑ **Cosine Similarity:**
 - Measures similarity *between books*.
 - *Based on TF-IDF vectors*.
- ❑ **Recommendation Function:**
 - Inputs a book title.
 - Outputs top N similar books.

Example

- ❑ **Input Book:** "The Lovely Bones: A Novel"
- ❑ **Recommendations:**
 1. "The Lovely Bones"
 2. "Lovely in Her Bones"
 3. "Bones"
 4. "Bare Bones: A Novel"

```
Recommendations for 'The Lovely Bones: A Novel':  
1. The Lovely Bones  
2. Lovely in Her Bones  
3. Bones  
4. Bare Bones : A Novel  
5. Bare Bones : A Novel
```


Collaborative Filtering (SVD)

- ❑ Predicts user preferences based on past interactions.
- ❑ Utilizes patterns in user-item interactions.

SVD Approach

- ❑ **Data Preparation:**
 - Used the Surprise library.
 - Loaded data suitable for SVD.
- ❑ **Model Training:**
 - Trained SVD with 100 latent factors.
 - Over 30 epochs.
- ❑ **Evaluation:**
 - Tested on the test set.
 - Computed RMSE

SVD Performance

SVD RMSE: 1.5828

Collaborative Filtering with SVD

We implement collaborative filtering using **Singular Value Decomposition (SVD)**:

Data Preparation

Use the surprise library to load data suitable for SVD.

Model Training

Train the SVD model on the training set.

Evaluation

Test the model on the test set and compute RMSE to assess prediction accuracy.

Model Performance

The SVD model was trained using the Surprise library. Evaluation results:

- **RMSE: 1.5850**

Neural Collaborative Filtering (NCF)

- Uses neural networks to model complex interactions.
- Captures non-linear relationships.

NCF Performance

- **Test Loss (MSE): 2.5271**
- **Test MAE: 1.2402**
- **NCF RMSE: 1.5897**

NCF Approach

- **Data Preprocessing:**
 - Encoded user and item IDs.
 - Split data into training and testing sets.
- **Model Architecture:**
 - Embedding layers for users and items.
 - Dense layers to capture interactions.
 - Output layer predicts the rating.
- **Training:**
 - Used early stopping to prevent overfitting.

```
Epoch 1/20
151/151 ██████████ 5s 14ms/step - loss: 40.3619 - mae: 5.6047 - val_loss: 2.8695 - val_mae: 1.3587
Epoch 2/20
151/151 ██████████ 2s 8ms/step - loss: 4.6134 - mae: 1.7183 - val_loss: 2.6767 - val_mae: 1.3056
Epoch 3/20
151/151 ██████████ 1s 8ms/step - loss: 3.8251 - mae: 1.5514 - val_loss: 2.5743 - val_mae: 1.2608
Epoch 4/20
151/151 ██████████ 2s 8ms/step - loss: 3.4515 - mae: 1.4691 - val_loss: 2.5816 - val_mae: 1.2670
Epoch 5/20
151/151 ██████████ 1s 8ms/step - loss: 3.2146 - mae: 1.4206 - val_loss: 2.6077 - val_mae: 1.2734
Epoch 6/20
151/151 ██████████ 1s 8ms/step - loss: 3.0457 - mae: 1.3766 - val_loss: 2.6021 - val_mae: 1.2710
Epoch 6: early stopping
Restoring model weights from the end of the best epoch: 3.
602/602 ██████████ 1s 1ms/step - loss: 2.5322 - mae: 1.2487
NCF Test Loss (MSE): 2.5743191242218018
NCF Test Mean Absolute Error (MAE): 1.2607836723327637
```


Hybrid Recommendation System

- Combines collaborative and content-based filtering.
- Leverages strengths of both methods.

Hybrid Approach

- **Components:**
 - SVD Prediction (Weight: 0.7)
 - Content-Based Prediction (Weight: 0.2)
 - Global Mean Rating (Weight: 0.1)
- **Hybrid Score:**
 - Weighted sum of the components.

Hybrid Performance

- **Sample Hybrid Score: 6.3660**
- **Hybrid RMSE: 2.1550**

Hybrid Recommendation System

We create a hybrid recommendation system that combines:

Collaborative Filtering Prediction (SVD)

- Captures user-item interactions.

Content-Based Prediction

- Uses cosine similarity of book titles.

Global Mean Rating

- Acts as a baseline.



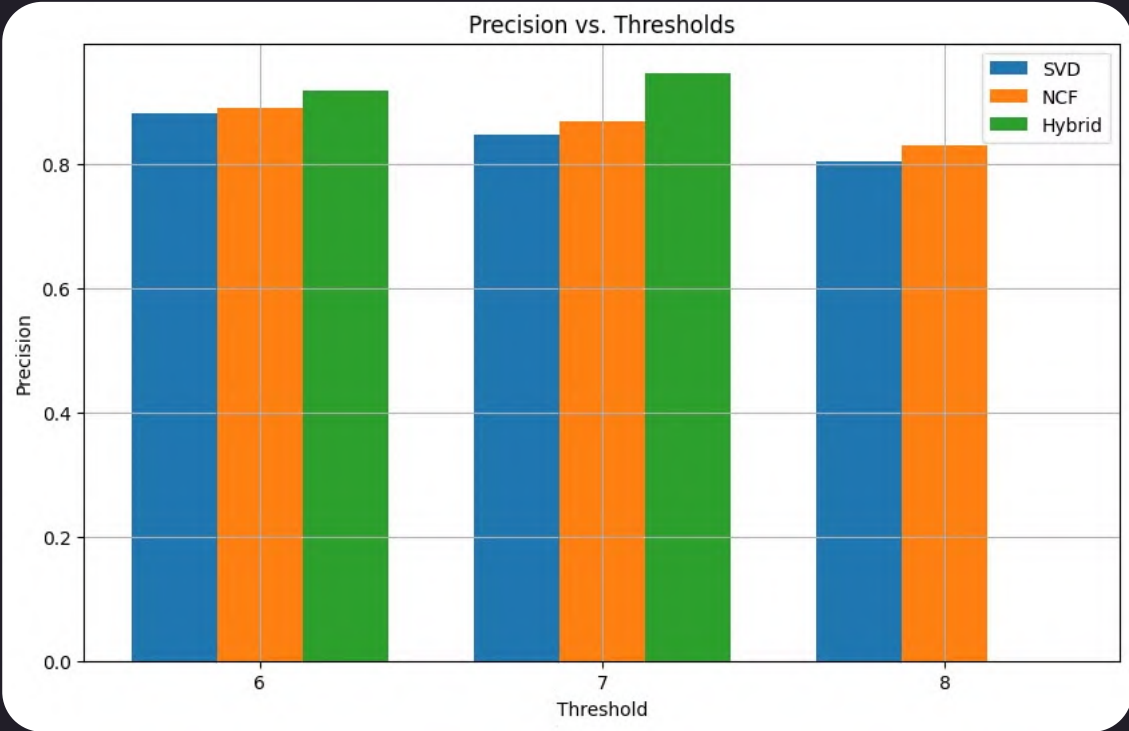
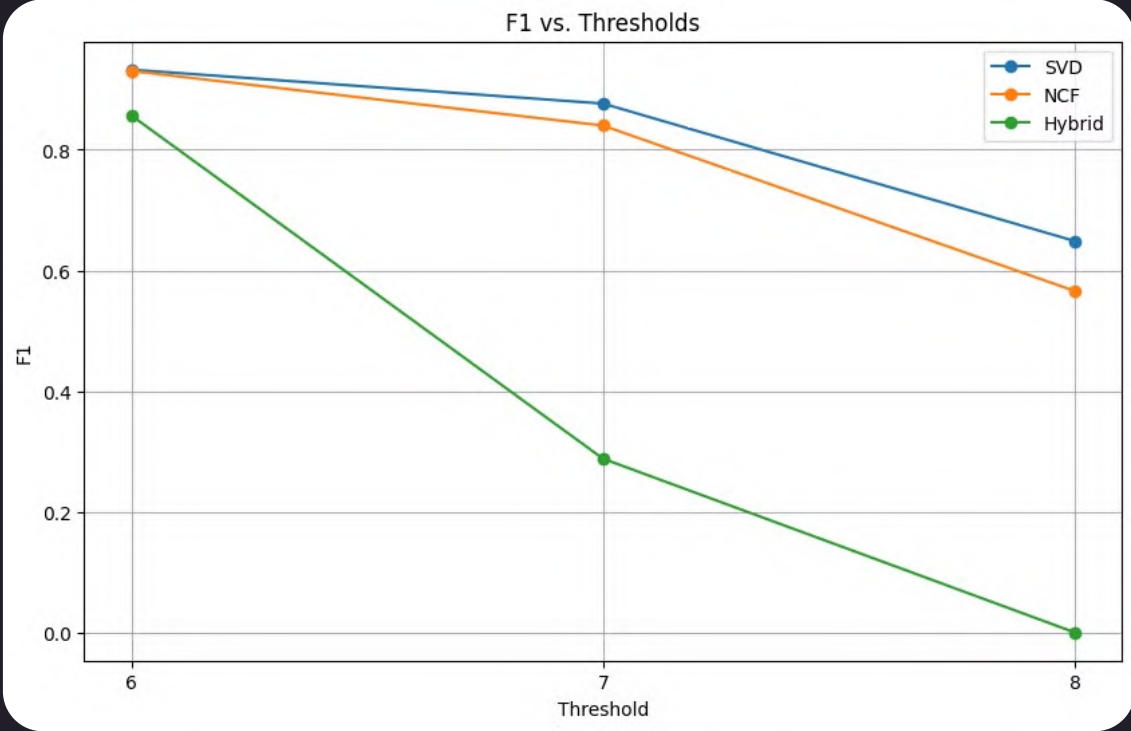

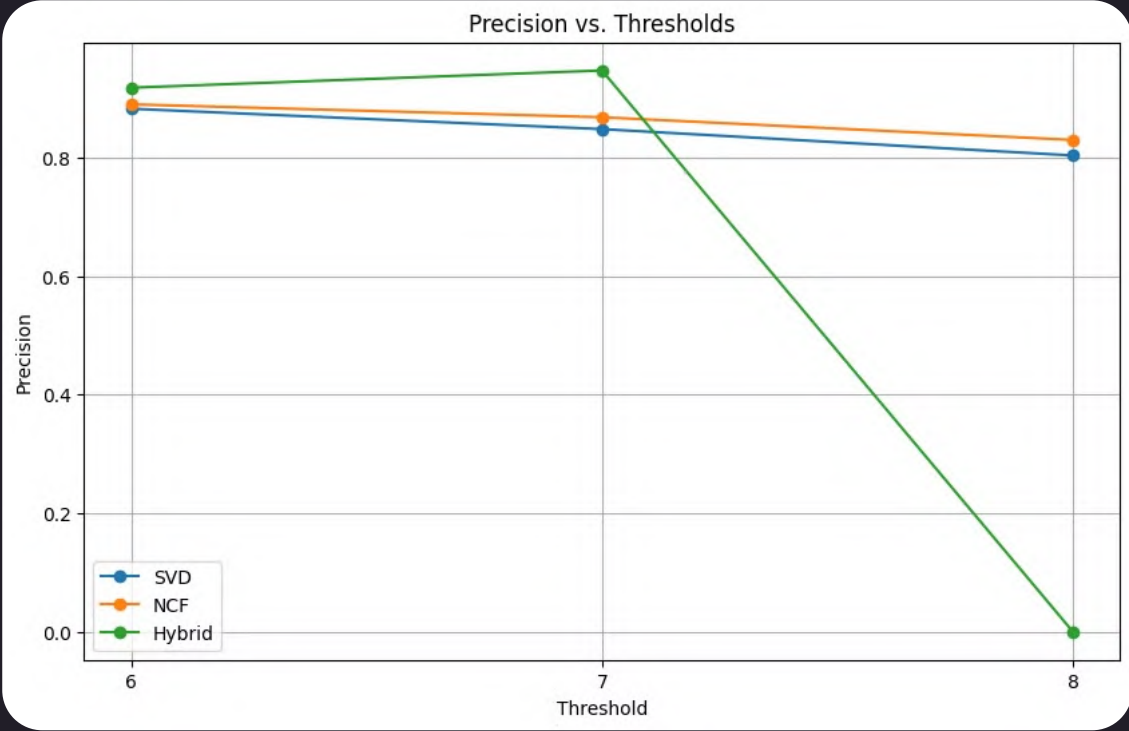
We assign weights to each component to control their influence on the final recommendation score.

Output

- **Hybrid Recommendation Score** for a sample: **6.2515**
- **Hybrid RMSE: 2.1578**

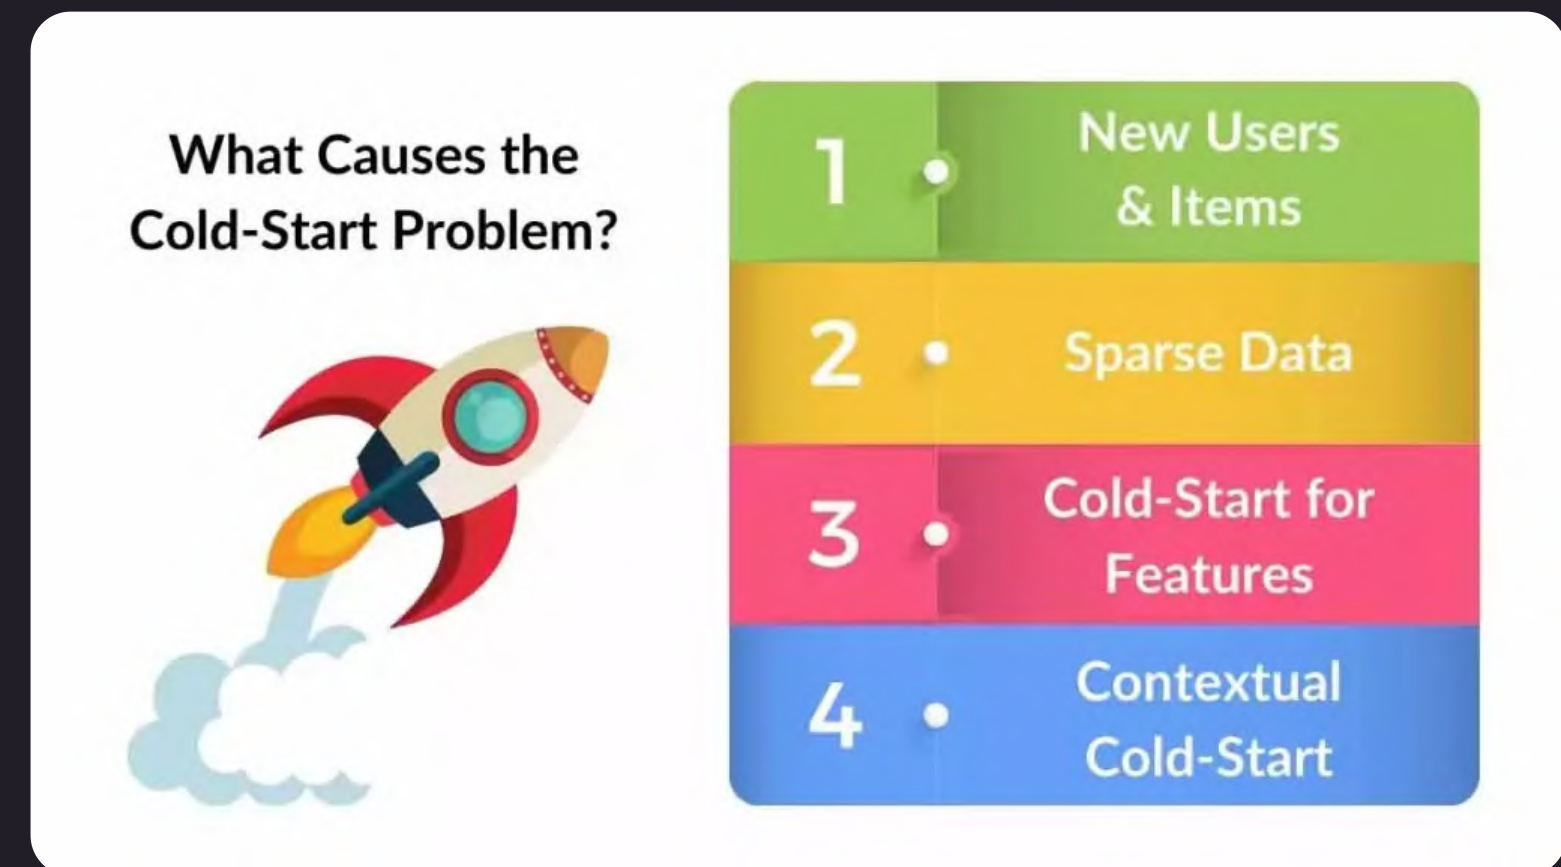

RMSE Comparison

Model	RMSE
SVD	1.5828
NCF	1.5897
Hybrid	2.1550
User Mean Baseline	1.6774
Item Mean Baseline	1.7476



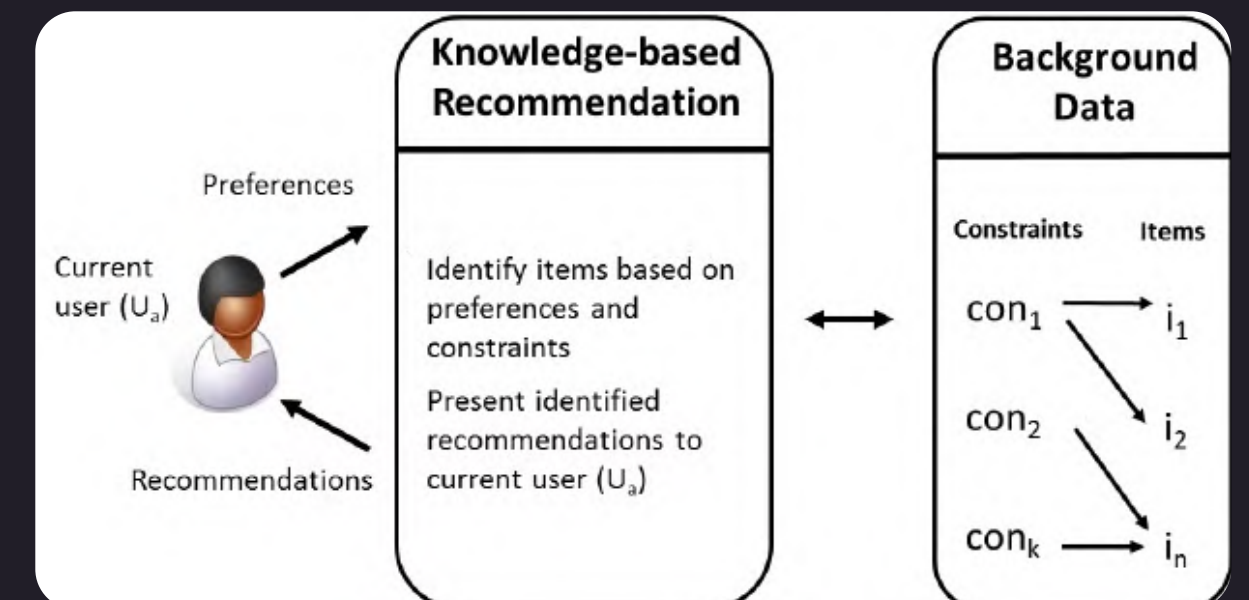


Challenges

- ❑ **Data Sparsity:**
 - Limited ratings per user/book.
 - Affects model learning.
- ❑ **Cold-Start Problem:**
 - New users/items lack historical data.
- ❑ **Computational Complexity:**
 - NCF requires more resources.
- ❑ **Evaluation Trade-Offs:**
 - Balancing precision and recall is crucial.



Future Work

- **Hyperparameter Tuning:**
 - Optimize NCF and hybrid models.
- **Hybrid Model Adjustment:**
 - Rebalance weights to improve recall.
- **Advanced Techniques:**
 - Explore other algorithms like KB recommendations.



Conclusion

Successfully implemented and evaluated various models.

SVD model was most effective.

NCF model showed promise.

Hybrid model needs adjustment.

The background of the slide features a dark, starry night sky with a gradient of deep blues and purples. In the lower portion, the silhouettes of sand dunes are visible, with a prominent, dark, triangular-shaped dune in the center-left foreground. The overall mood is serene and professional.

Thank You For Listening!

Now any questions?

References

- ❑ **Book-Crossing Dataset:**
[Kaggle](#)
- ❑ **Surprise Library**
Documentation: [Surprise](#)
- ❑ **TensorFlow Keras**
Documentation: [TensorFlow](#)
- ❑ **Scikit-learn Documentation:**
[Scikit-learn](#)
- ❑ **Pandas Documentation:**
[Pandas](#)



Want to make a presentation like this one?

Start with a fully customizable template, create a beautiful deck in minutes, then easily share it with anyone.

Create a presentation (It's free)